

Parallelisation of a Staggered Grid solver

Heisig, Hammer, Ernst

FAU

February 2, 2014

Outline

Parallelization basics

- ▶ Why parallelize your code?

Pro

- ▶ more compute power
- ▶ more memory
- ▶ parallel computing is the future

Con

- ▶ added code complexity
- ▶ communication overhead
- ▶ Increased power consumption

Don't parallelize without profiling and performance modelling!

- ▶ MPI in a nutshell

The **Message Passing Interface**

- ▶ call your program with `mpirun -np <N> <NAME> <ARGS>`
- ▶ spawns <N> identical processes
- ▶ only `MPI_MPI_Comm_rank(...)` gives different results

Typical usage:

Implementation

► Implementation

The following steps must be parallelized

- `SOR::solve()`
- `SOR::residual()`
- `SOR::normalize()`
- `determineNextDT()`
- `refreshBoundaries()`
- `computeFG()`
- `composeRHS()`
- `updateVelocities()`

Most of the time is spent in the SORSolver, so this is the focus.

► Domain partitioning

- Usually domain is split in roughly quadratic tiles
- We chose the simpler approach: Split in horizontal stripes
- Pro
 - easier to implement
 - fast access patterns along the cachelines
- Con
 - bad surface / size ratio for large number of processes

Results

- ▶ Results

Was it worth the effort?

Explanation:

- ▶ SOR or Jacobi solver does not scale well

Use a better algorithm before writing parallel code!

Possible improvements for numerical codes

- ▶ Possible improvements for numerical codes
Use LISP