

Parrallelisation of a Staggered Grid solver

Heisig, Hammer, Ernst

February 2, 2014

Outline

- 1 Parallelization basics
- 2 Implementation
- 3 Results
- 4 Possible improvements for numerical codes

Why parallelize your code?

Pro

- more compute power
- more memory
- parallel computing is the future

Con

- added code complexity
- communication overhead
- Increased power consumption

Don't parallelize without profiling and performance modelling!

The Message Passing Interface

- call your program with `mpirun -np <N> <NAME> <ARGS>`
- spawns <N> identical processes
- only `MPI_MPI_Comm_rank(...)` gives different results

Typical usage:

- split domain between all processes
- perform local updates
- exchange the borders
- repeat

The following steps must be parallelized

- `SOR::solve()`
- `SOR::residual()`
- `SOR::normalize()`
- `determineNextDT()`
- `refreshBoundaries()`
- `computeFG()`
- `composeRHS()`
- `updateVelocities()`

Most of the time is spent in the `SORSolver`, so this is the focus.

Domain partitioning

- Usually domain is split in roughly quadratic tiles
- We chose the simpler approach: Split in horizontal stripes

Pro

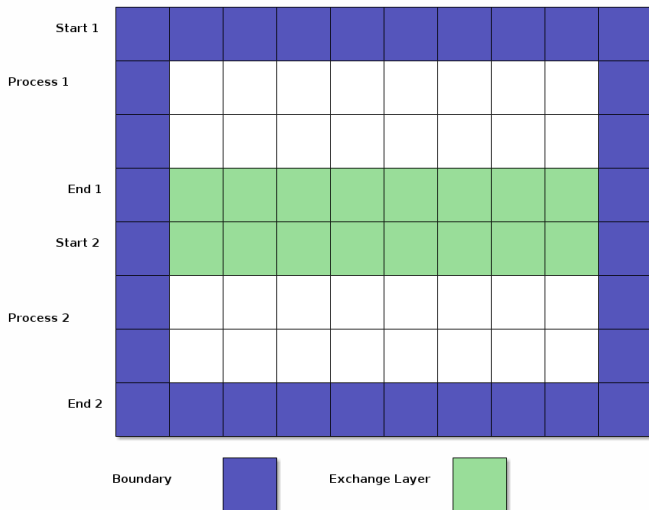
- easier to implement
- fast access patterns along the cachelines

Con

- bad surface / size ratio for large number of processes
- more communication overhead

Domain partitioning (cont.)

Example of a 8 x 9 domain with 2 processes



Continuous migration

How do we migrate our serial codebase to a parallel one without the agonizing pain™ ?

Migration phase:

- Every process still has all the data
- Parallelize only one operation at a time
- Methods can be tested individually

When all methods are converted, switch the Array implementation to store only local elements.

Was it worth the effort?

Explanation:

- SOR or Jacobi solver does not scale well

Use a better algorithm before writing parallel code!

Possible improvements for numerical codes

Use LISP